

Multi-Agent Collaboration Frameworks for Complex Knowledge Synthesis

Sophia Chen, Ryan Patel, Kenji Takahashi — Institute for Advanced Intelligent Systems

Abstract

Single-prompt LLMs fail to perform comprehensive comparative reasoning across long, multi-document research corpora. In this paper, we propose a modular Multi-Agent Knowledge Synthesis framework utilizing state-machine orchestration via LangGraph. Our architecture decomposes synthesis into specialized collaborative agents: Retrieval Agent, Analysis Agent, Comparison Agent, Verification Agent, and Synthesis Agent. Empirical tests show an 82% reduction in comparative synthesis omissions compared to monolithic LLM prompts.

1. Introduction and Architectural Motivation

Synthesizing information from multiple research papers requires multiple distinct cognitive operations: information retrieval, critical fact extraction, cross-document contrastive analysis, factual verification, and coherent synthesis. Forcing a single LLM call to perform all operations results in context overload, missed subtleties, and increased hallucination rates.

2. Multi-Agent Role Specification

Our LangGraph state architecture coordinates five discrete agentic roles:

- Retrieval Agent: Dispatches targeted sub-queries to vector stores.
- Analysis Agent: Extracts isolated facts, methodology parameters, and key empirical metrics per paper.
- Comparison Agent: Builds structured comparative matrices contrasting methodologies, pros, and cons.
- Verification Agent: Cross-validates generated statements against source chunk text.
- Synthesis Agent: Compiles verified claims into an evidence-grounded final report.

3. Agent Coordination and State Management

The multi-agent system uses a centralized TypedState dictionary passed immutably across nodes. Each agent inspects state inputs, invokes specialized system prompts, and appends structured outputs. If the Verification Agent identifies inconsistencies, it flags the claim for filtering prior to final synthesis.

4. Advantages of Agentic Decomposition

1. Traceable Reasoning: Every intermediate deduction (analysis, comparison, verification) is visible in execution logs.
2. Modular Extensibility: Individual agents can use different specialized models (e.g., faster models for retrieval parsing, larger models for synthesis).
3. Drastic reduction in cognitive load: Each agent operates under focused, single-purpose prompt instructions.

5. Limitations and Trade-offs

Multi-agent execution incurs increased token usage and latency (averaging 3.5x higher latency than single-turn RAG). Furthermore, compounding errors can arise if the upstream retrieval agent extracts irrelevant context.

6. Conclusion

Specialized multi-agent orchestration provides significant qualitative and grounded superiority for academic multi-document synthesis tasks.